

EECS151/251A

Introduction to Digital Design and ICs


Lecture 10: RISC-V Pipelining Sophia Shao



Intel's plan to license x86 cores for chips with Arm, RISC-V and more inside

Intel will license its most important asset, the x86 architecture, to those who want to make custom silicon. Depending on the application, customers will be able to mix up x86, Arm and RISC-V CPU cores as well as hardware acceleration units in a custom-designed chip that Intel fabricates.

"We have what we call a multi ISA strategy. That's the first time in Intel's history we'll license x86 soft cores and hard cores to customers who would like to develop chips," Bob Brennan, vice president of customer solutions engineering at Intel's Foundry Services, told *The Register*.

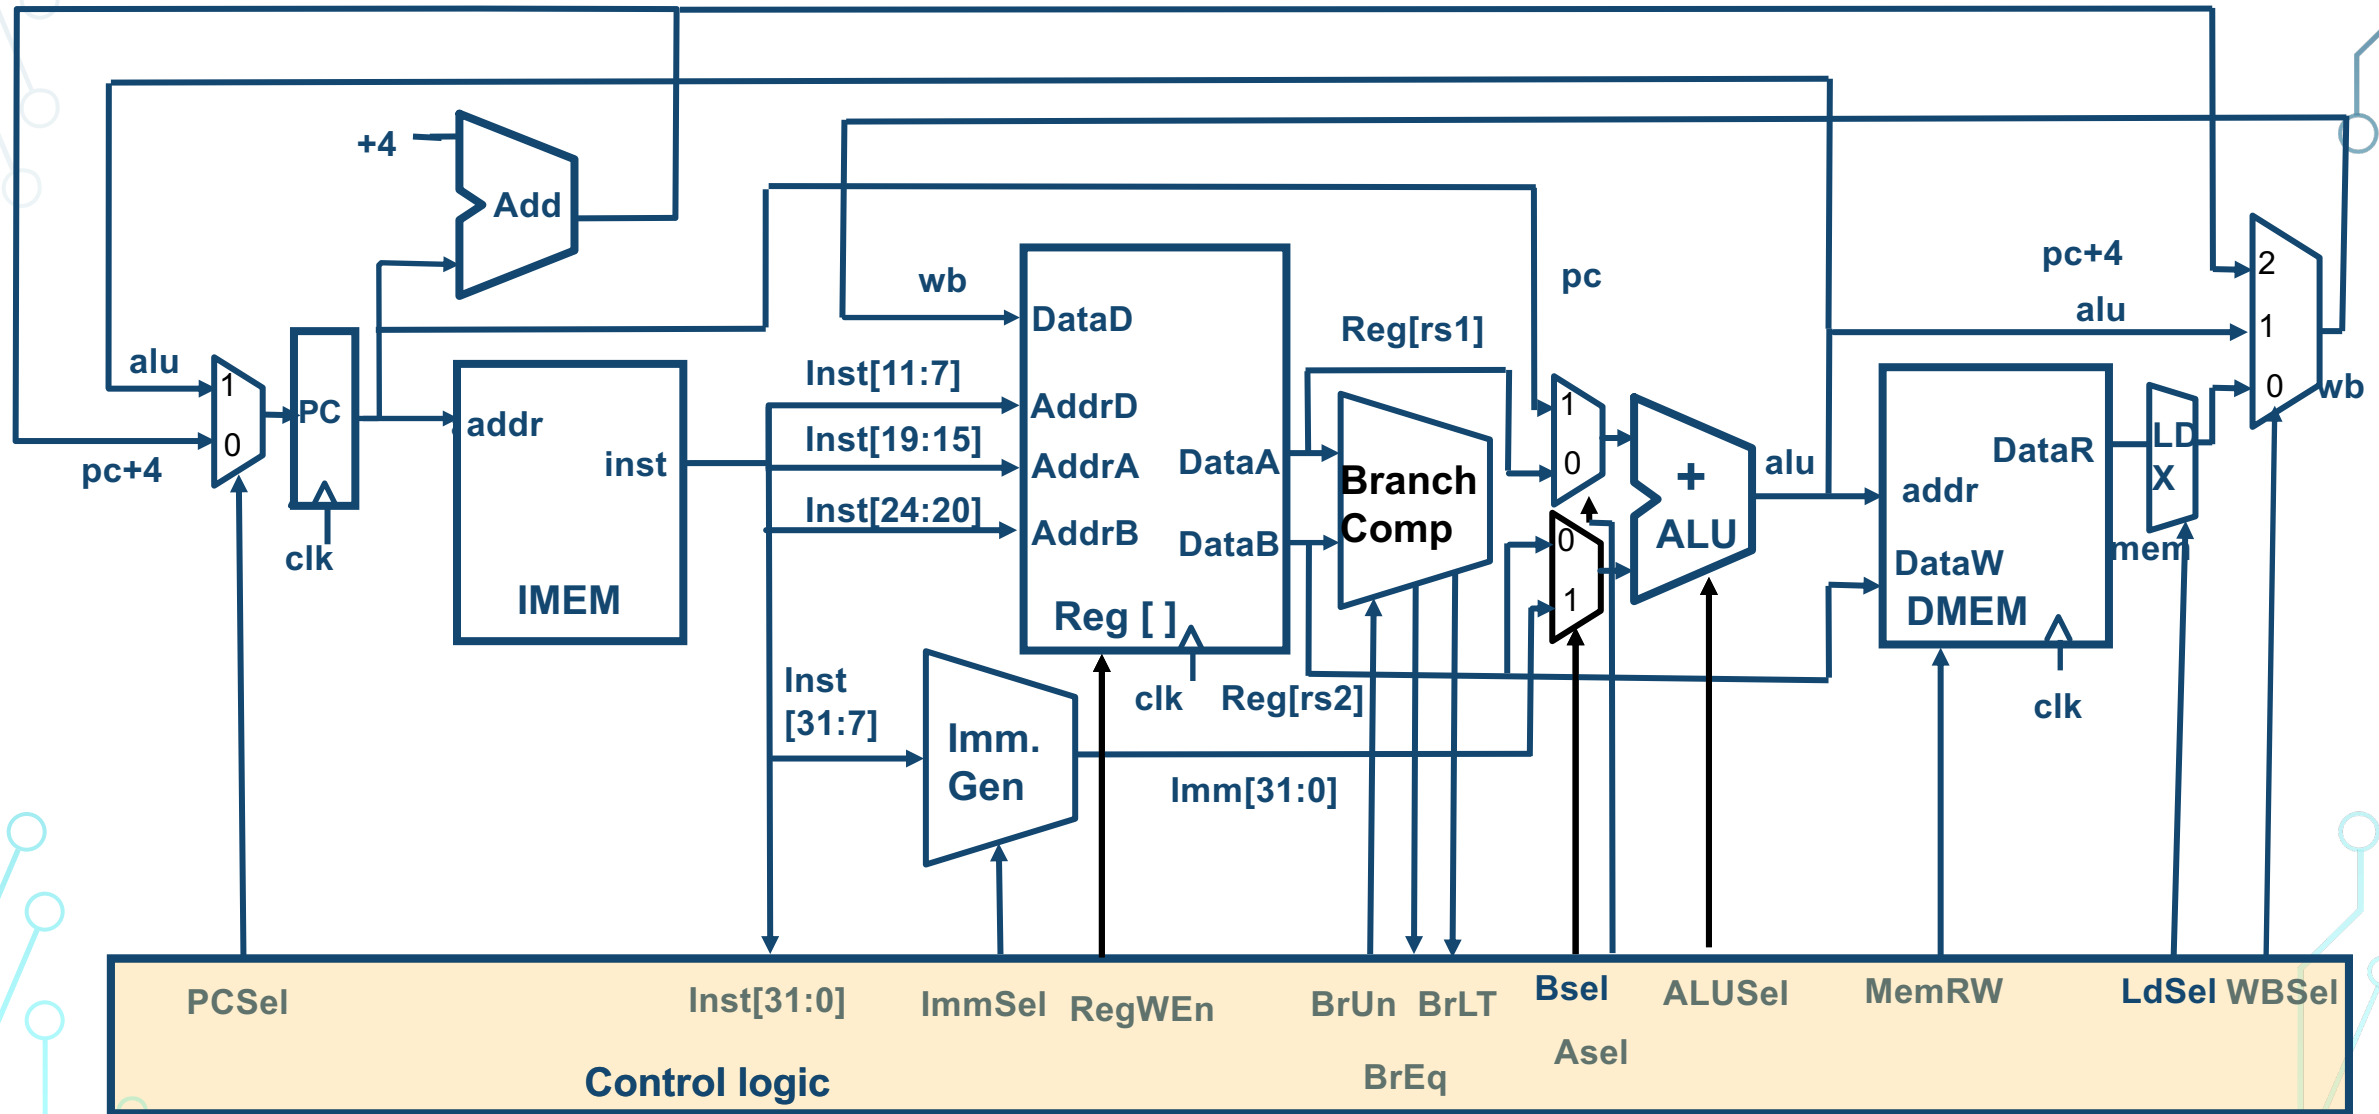


https://www.theregister.com/2022/02/14/intel_x86_licensing/

Review

- We have covered the implementation of the base ISA for RV32I!!!
 - Get yourself familiar with the ISA Spec.
- Instruction type:
 - R-type
 - I-type
 - S-type
 - B-type
 - J-type
 - U-type
- Implementation suggested is straightforward, yet there are modalities in how to implement it well at gate level.
- Single-cycle datapath is slow – need to pipeline it

Complete RV32I Datapath with Control

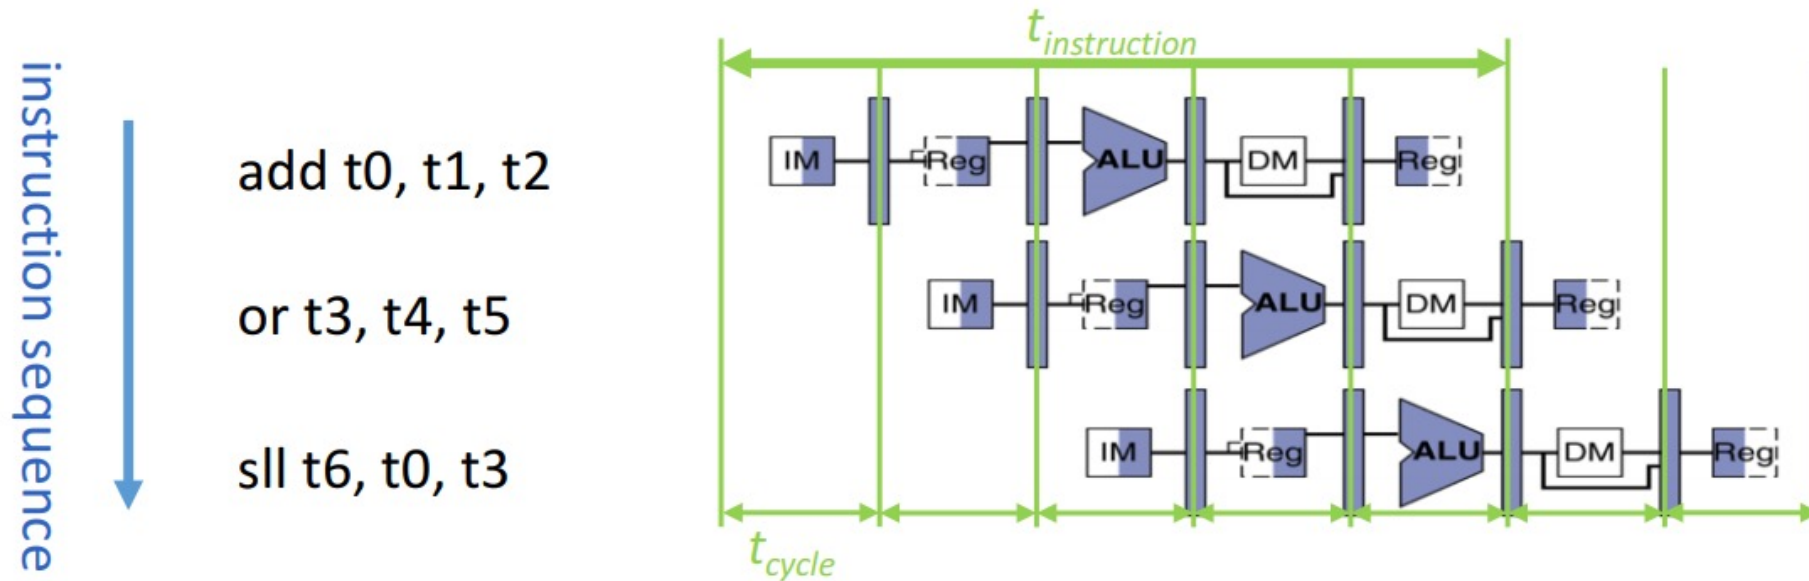




- **RISC-V Pipelining**

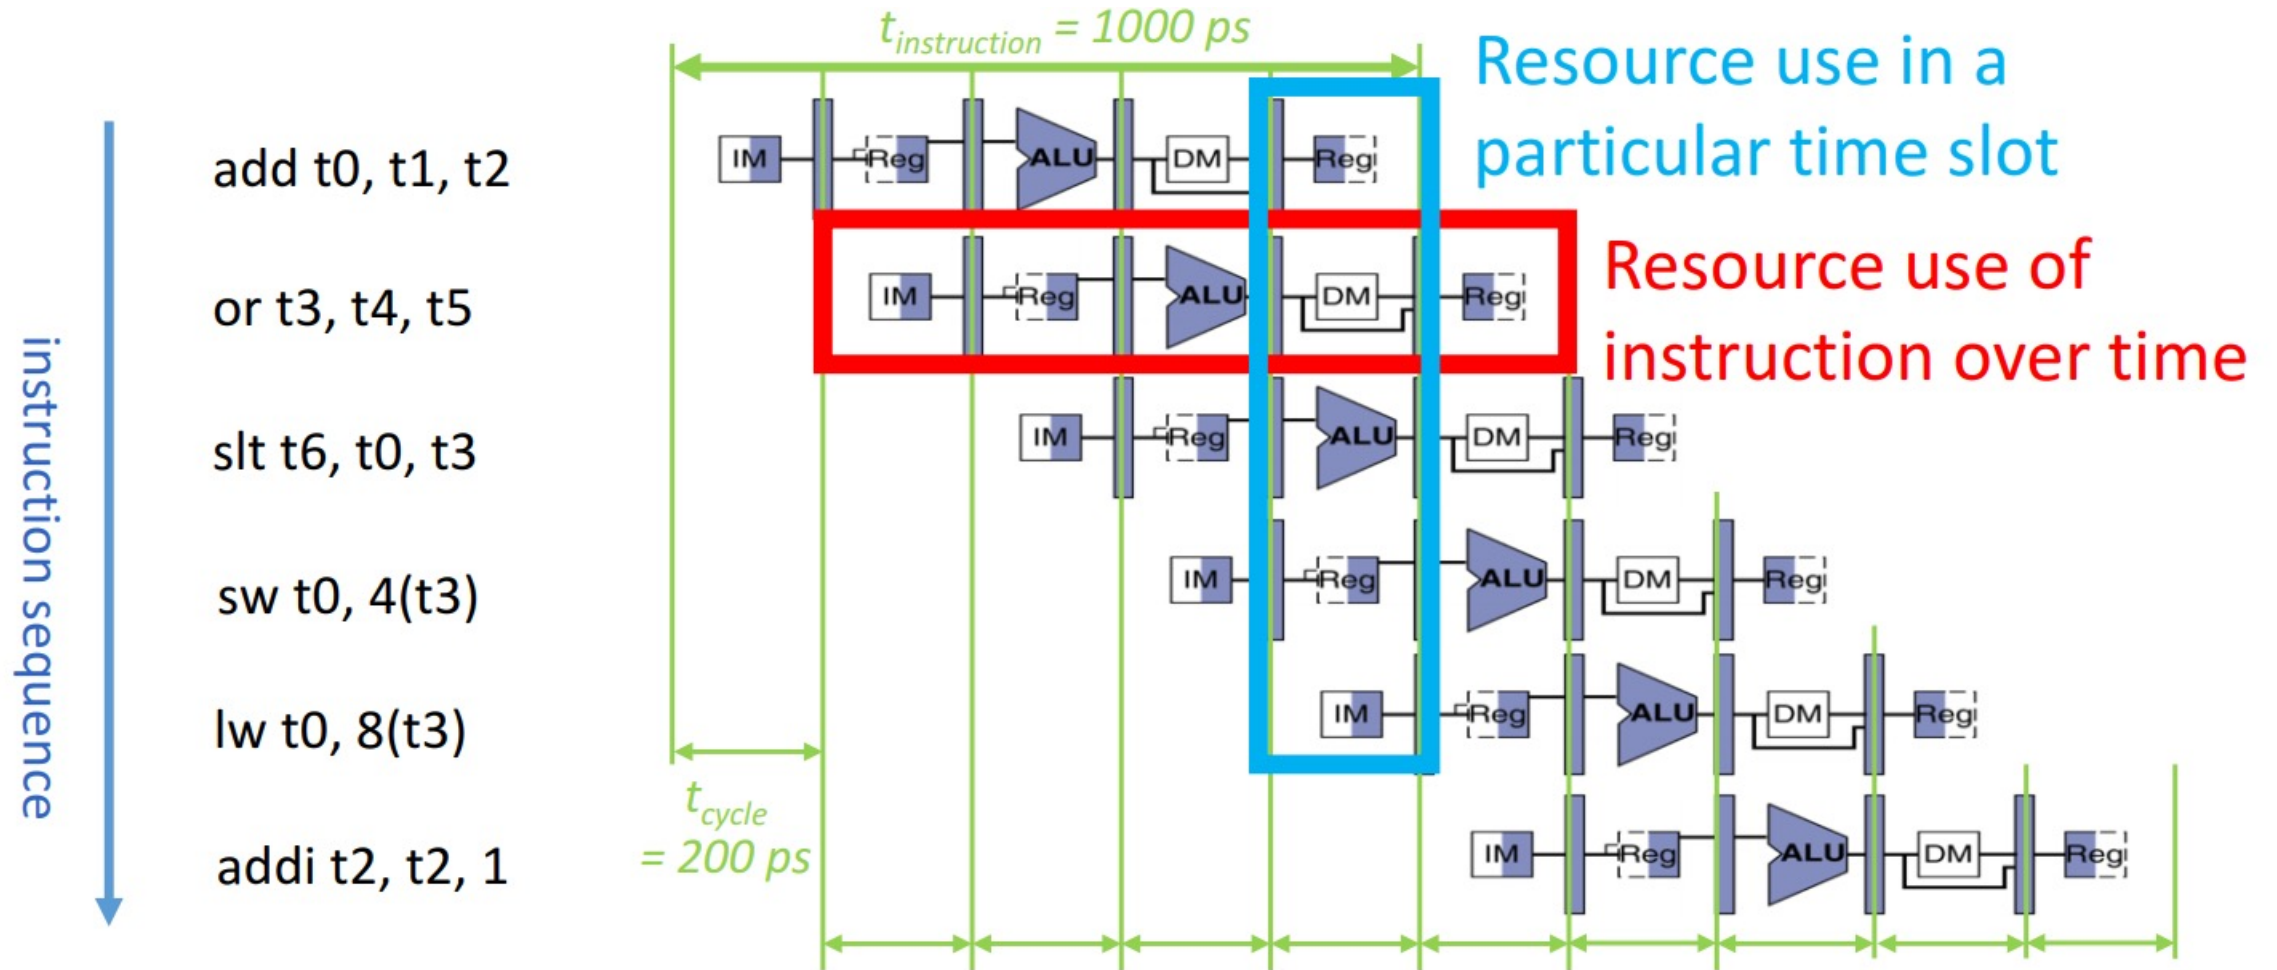
- **5-Stage Pipeline**
- **Pipeline Hazards**
 - **Structural**
 - **Data**
 - **Control**

Pipelining with RISC-V

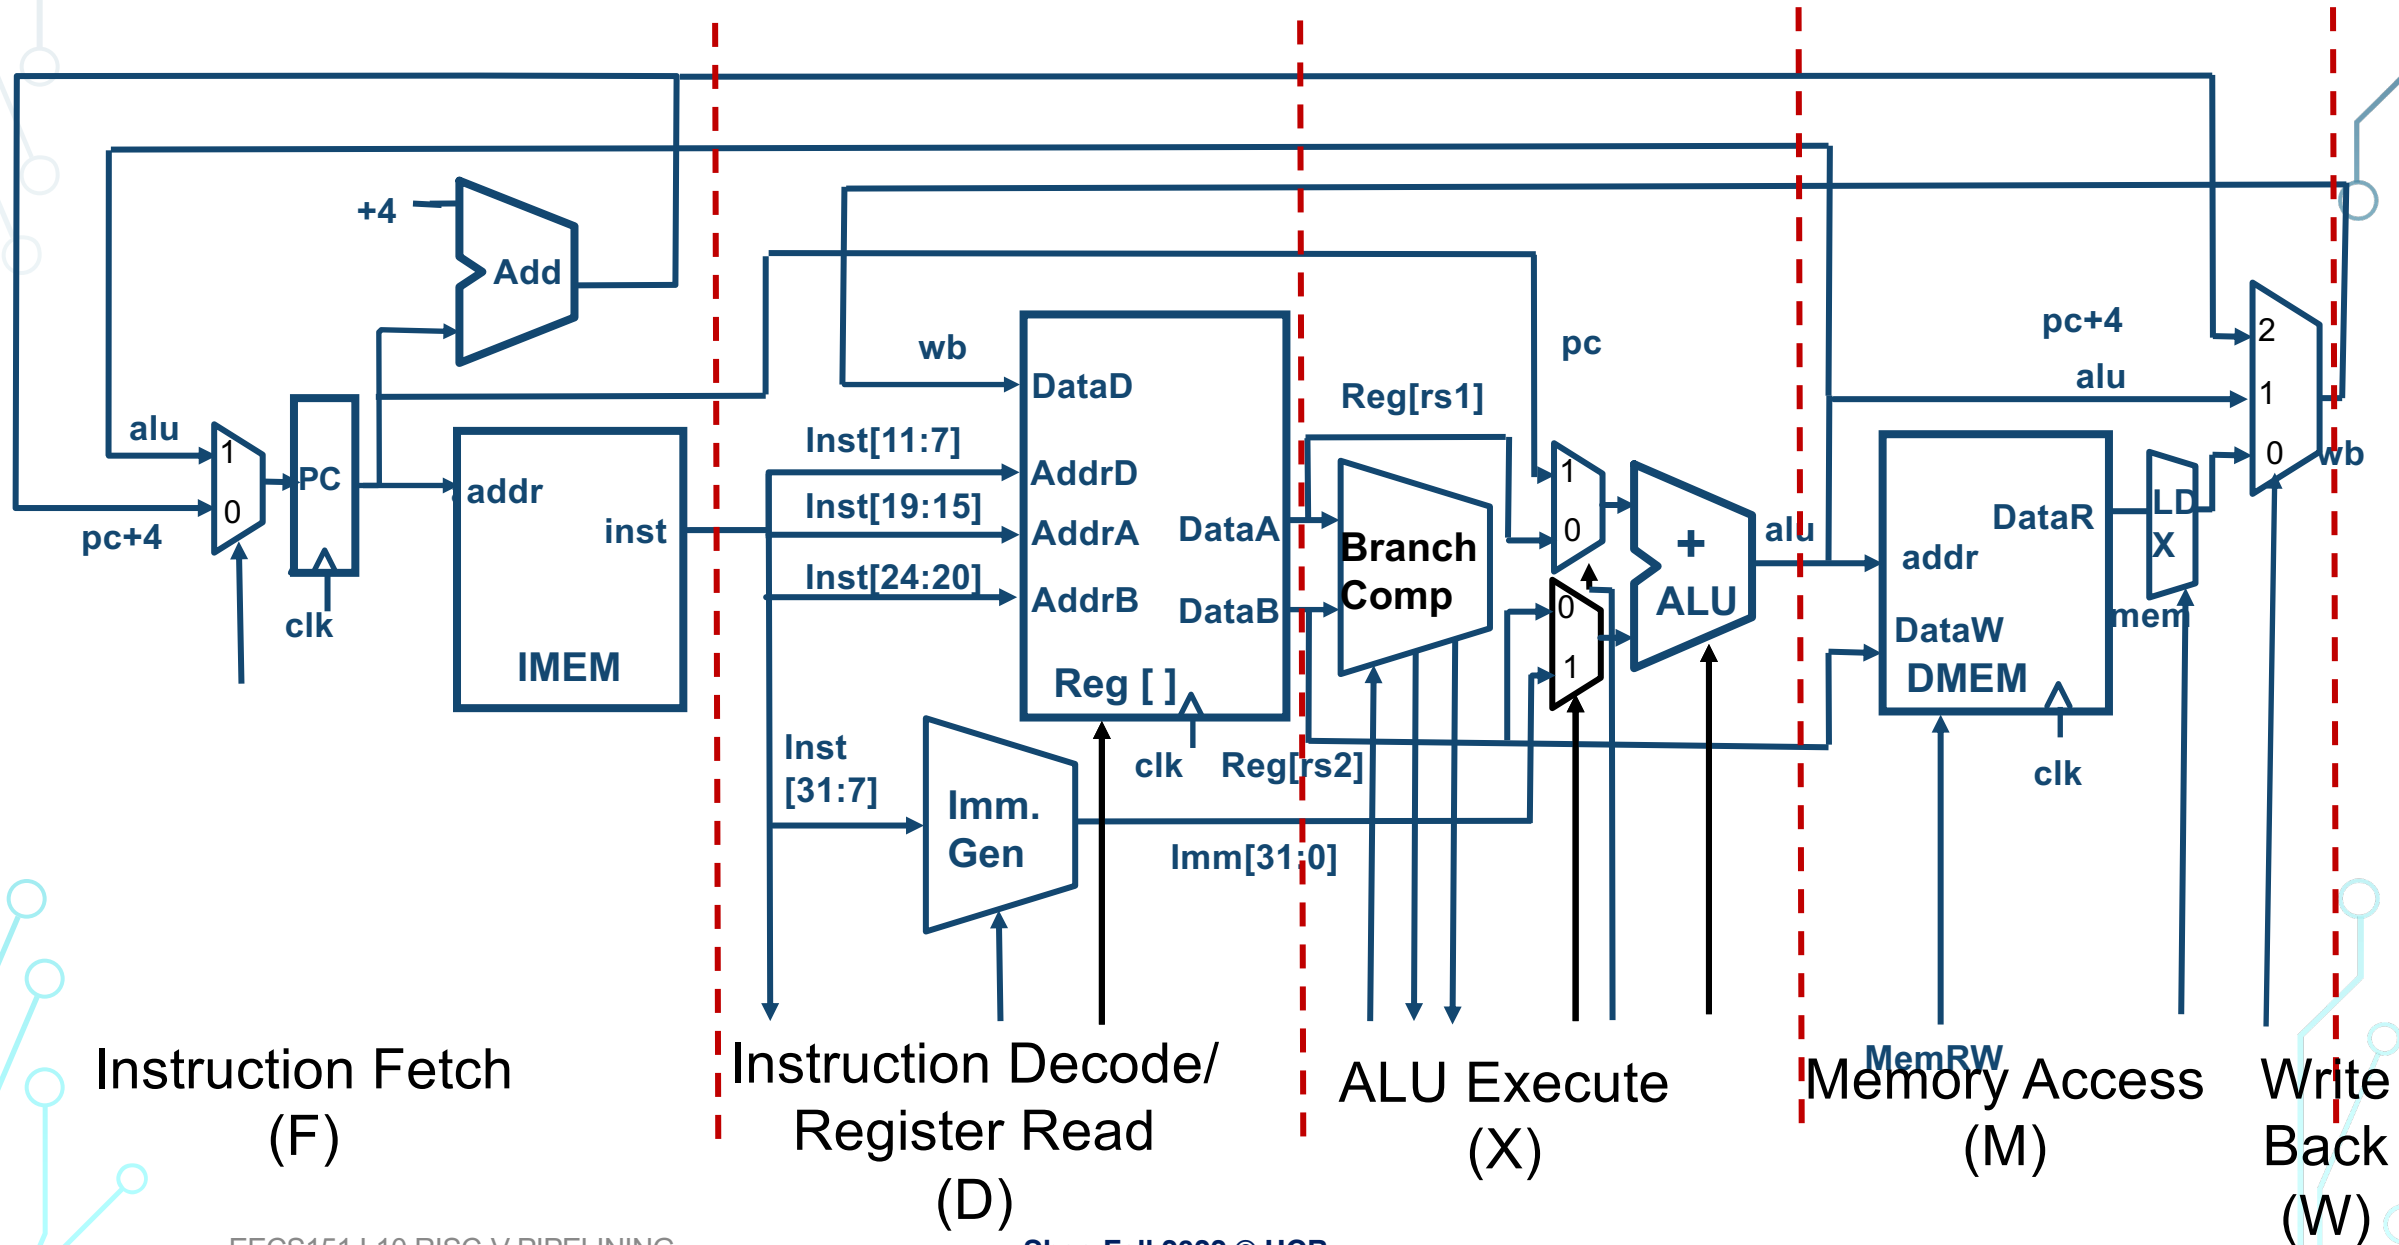


	Single Cycle	Pipelining
Timing	$t_{step} = 100 \dots 200 \text{ ps}$	$t_{cycle} = 200 \text{ ps}$
	Register access only 100 ps	All cycles same length
Instruction time, $t_{instruction}$	$= t_{cycle} = 800 \text{ ps}$	1000 ps
Clock rate, f_s	$1/800 \text{ ps} = 1.25 \text{ GHz}$	$1/200 \text{ ps} = 5 \text{ GHz}$
Relative speed	1 x	4 x


Pipelining with RISC-V



Complete RV32I Datapath with Control




Pipelining RV32I Datapath




Must pipeline instruction along with data, so control operates correctly in each stage

Different Instructions in Flight



Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation
 - Information is stored in pipeline registers for use by later stages





- **RISC-V Pipelining**

- **5-Stage Pipeline**
- **Pipeline Hazards**
 - **Structural**
 - **Data**
 - **Control**

Pipelining Hazards

A *hazard* is a situation that prevents starting the next instruction in the next clock cycle

1) *Structural hazard*

- A required resource is busy (e.g. needed in multiple stages)

2) *Data hazard*

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

3) *Control hazard*

- Flow of execution depends on previous instruction

Structural Hazard

- **Problem:** Two or more instructions in the pipeline compete for access to a single physical resource
- **Solution 1:** Instructions take it in turns to use resource, some instructions have to stall
- **Solution 2:** Add more hardware to machine
- *Can always solve a structural hazard by adding more hardware*

Structural Hazards I: Regfile

- Each instruction:
 - can read up to two operands in decode stage
 - can write one value in writeback stage
- Avoid structural hazard by having separate “ports”
 - two independent read ports and one independent write port
- Three accesses per cycle can happen simultaneously
 - We will see a circuit design for multi-ported register files later in the class
- Processors with multiple execution paths have larger number of ports

Structural Hazard II : Memory Access

instruction sequence


add t0, t1, t2

or t3, t4, t5

slt t6, t0, t3

sw t0, 4(t3)

lw t0, 8(t3)



- Instruction and data memory used simultaneously
✓ Use two separate memories

Structural Hazards – Summary

- Conflict for use of a resource
- In RISC-V pipeline with a single memory
 - Load/store requires data access
 - Without separate memories, instruction fetch would have to *stall* for that cycle
 - All other operations in pipeline would have to wait
- Pipelined datapaths require separate instruction/data memories
 - Or separate instruction/data caches
- RISC ISAs (including RISC-V) designed to avoid structural hazards
 - e.g. at most one memory access/instruction



- **RISC-V Pipelining**

- **5-Stage Pipeline**
- **Pipeline Hazards**
 - **Structural**
 - **Data**
 - **Control**

Data Hazard: Register Access

- Separate ports, but what if write to same value as read?
- Does **sw** in the example fetch the old or new value?

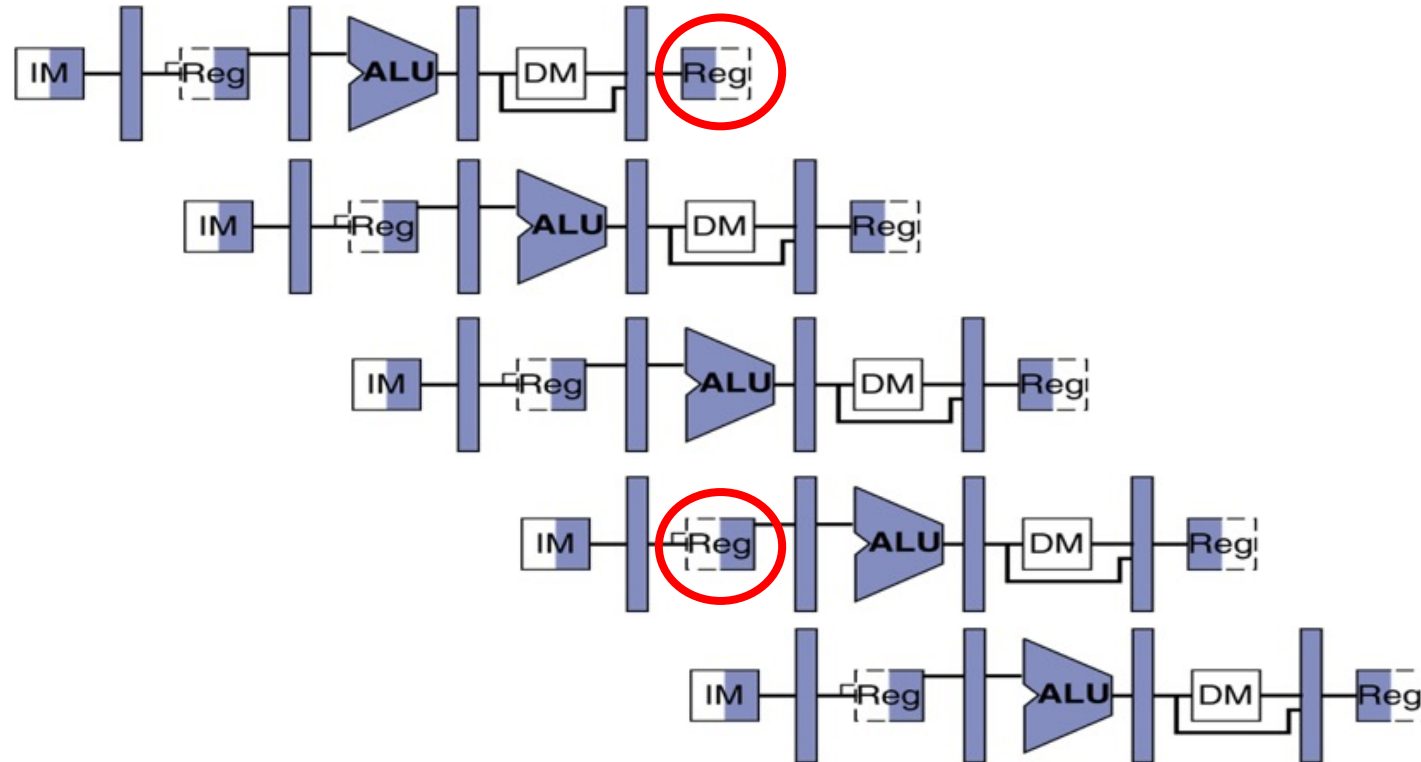
add t0, t1, t2

or t3, t4, t5

slt t6, t0, t3

sw t0, 4(t3)

lw t0, 8(t3)



instruction sequence

Register Access Policy

instruction sequence

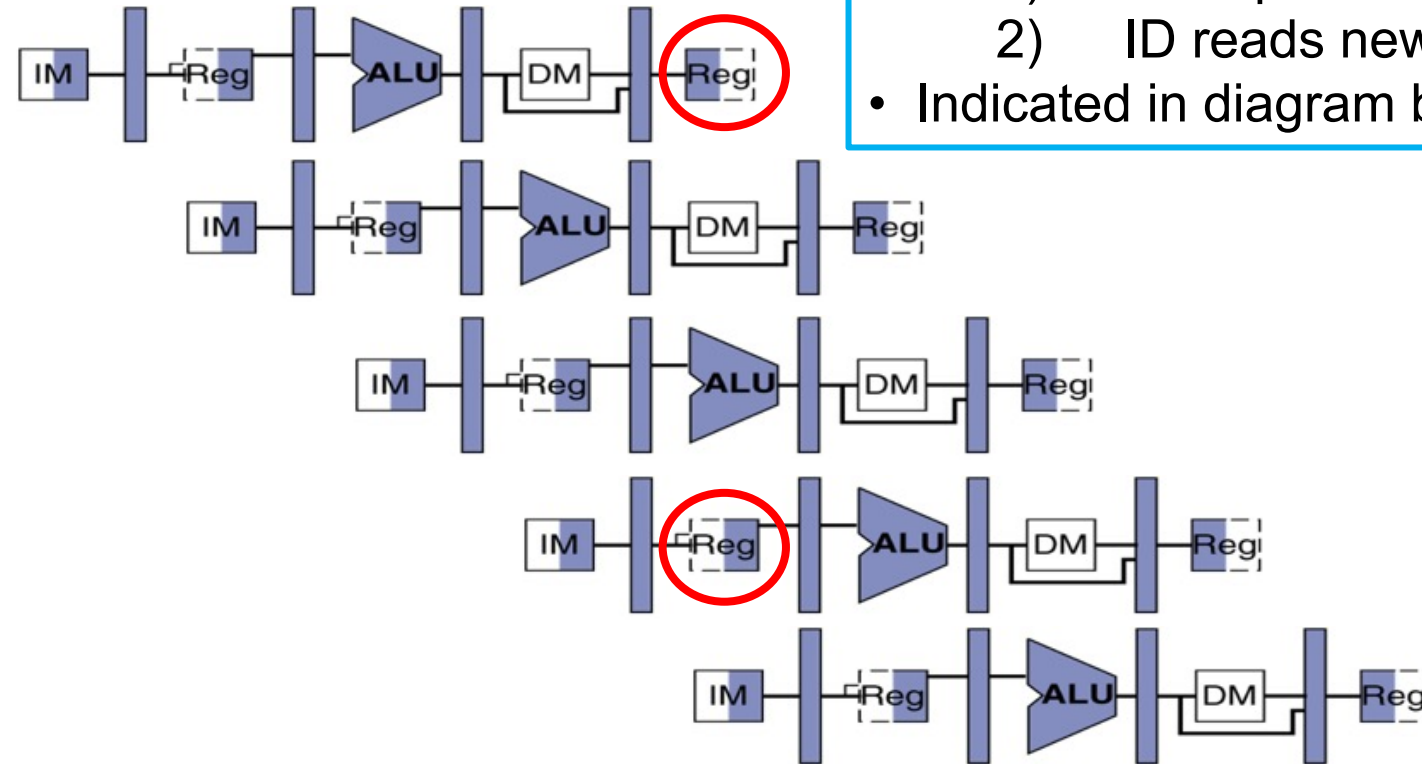
add t0, t1, t2

or t3, t4, t5

slt t6, t0, t3

sw t0, 4(t3)

lw t0, 8(t3)



- Exploit high speed of register file (100 ps)
 - 1) WB updates value
 - 2) ID reads new value
- Indicated in diagram by shading

Might not always be possible to write then read in same cycle, especially in high-frequency designs. Check assumptions in any question.

Data Hazard: ALU Result

Value of **s0**

5	5	5	5	5/9	9	9	9	9
---	---	---	---	-----	---	---	---	---

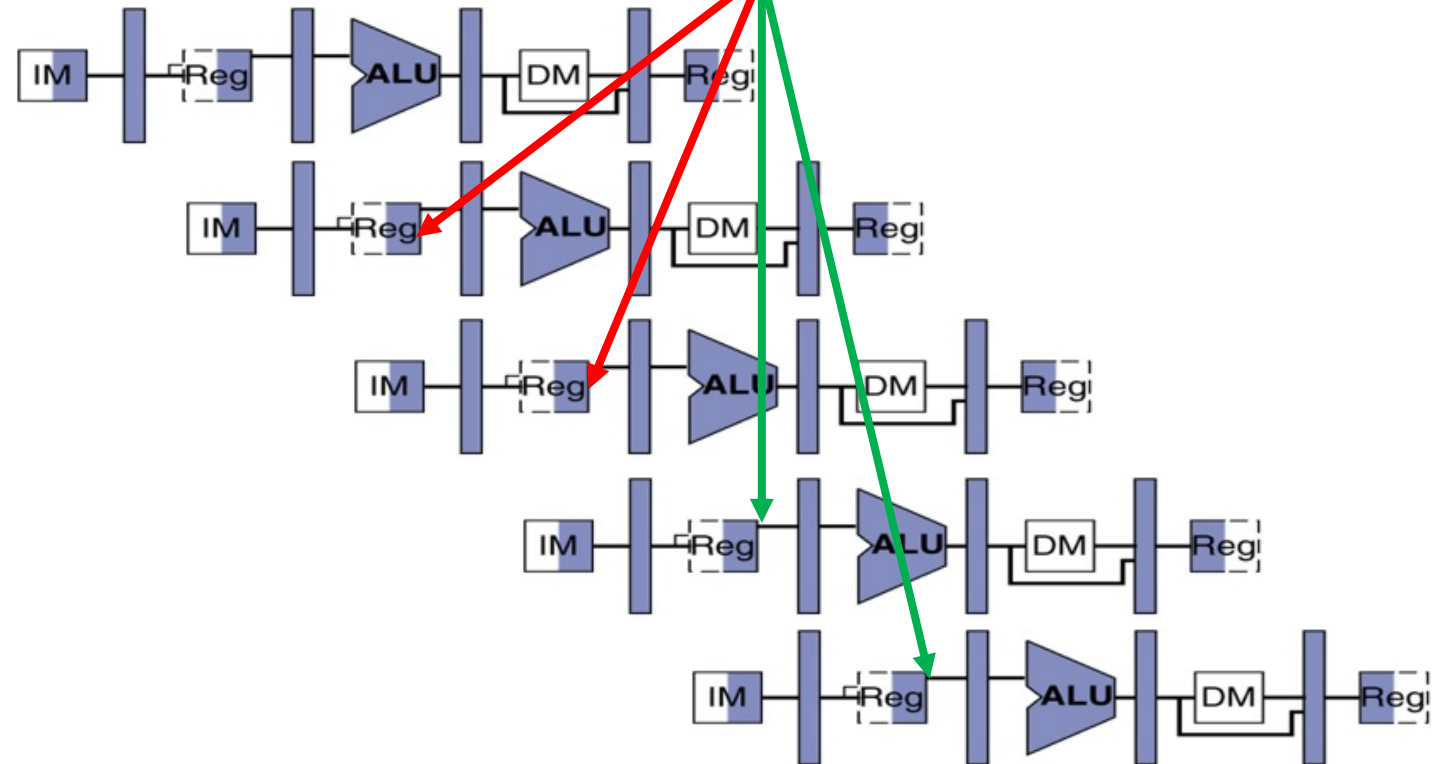
add **s0**, t0, t1

sub t2, **s0**, t0

or t6, **s0**, t3

xor t5, t1, **s0**

sw **s0**, 8(t3)




Without some fix, **sub** and **or** will calculate wrong result!

Solution 1: Stalling

- Problem: Instruction depends on result from previous instruction

- add s0, t0, t1
- sub t2, s0, t3




- Bubble:
 - effectively NOP: affected pipeline stages do “nothing”

Stalls and Performance


- Stalls reduce performance
 - But stalls are required to get correct results
- Compiler can arrange code or insert NOPs (writes to register x0) to avoid hazards and stalls
 - But requires knowledge of the pipeline structure

Solution 2: Forwarding

- Use result when it is computed
 - Don't wait for it to be stored in a register
 - Requires extra connections in the datapath



Solution 2: Forwarding



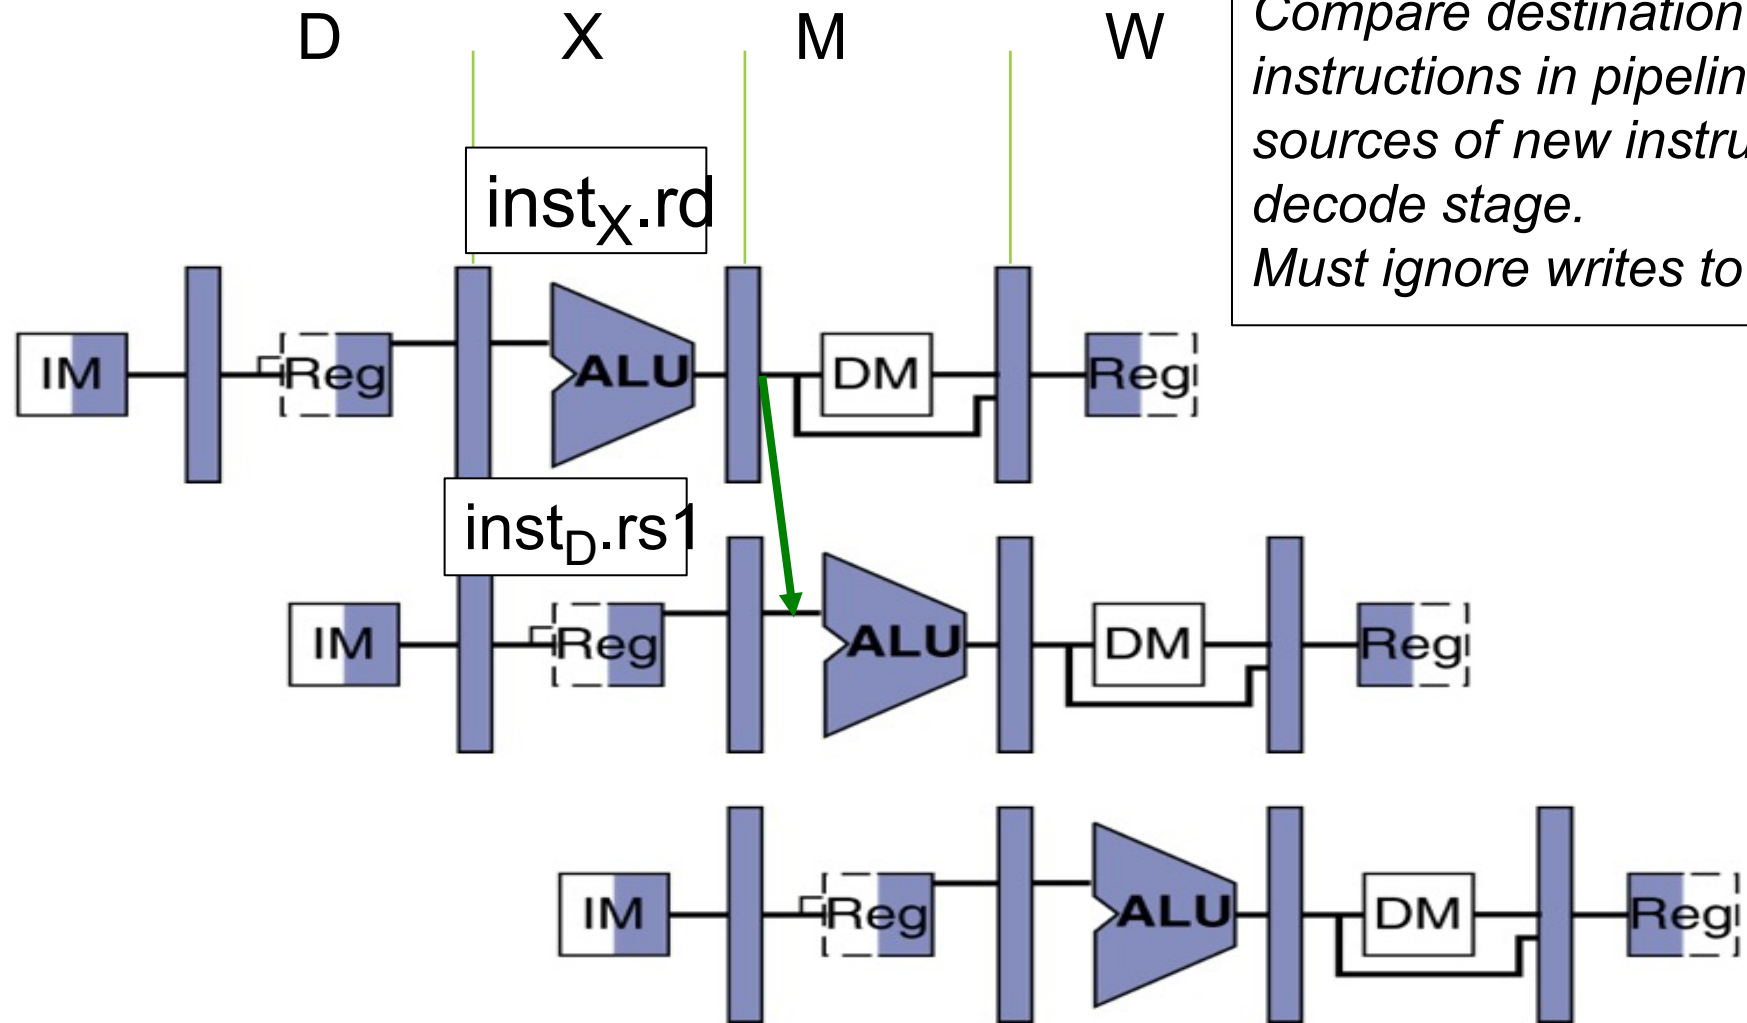
Forwarding: grab operand from pipeline stage,
rather than register file

Detect Need for Forwarding (example)

add **t0**, t1, t2

or t3, **t0**, t5

sub t6, **t0**, t3



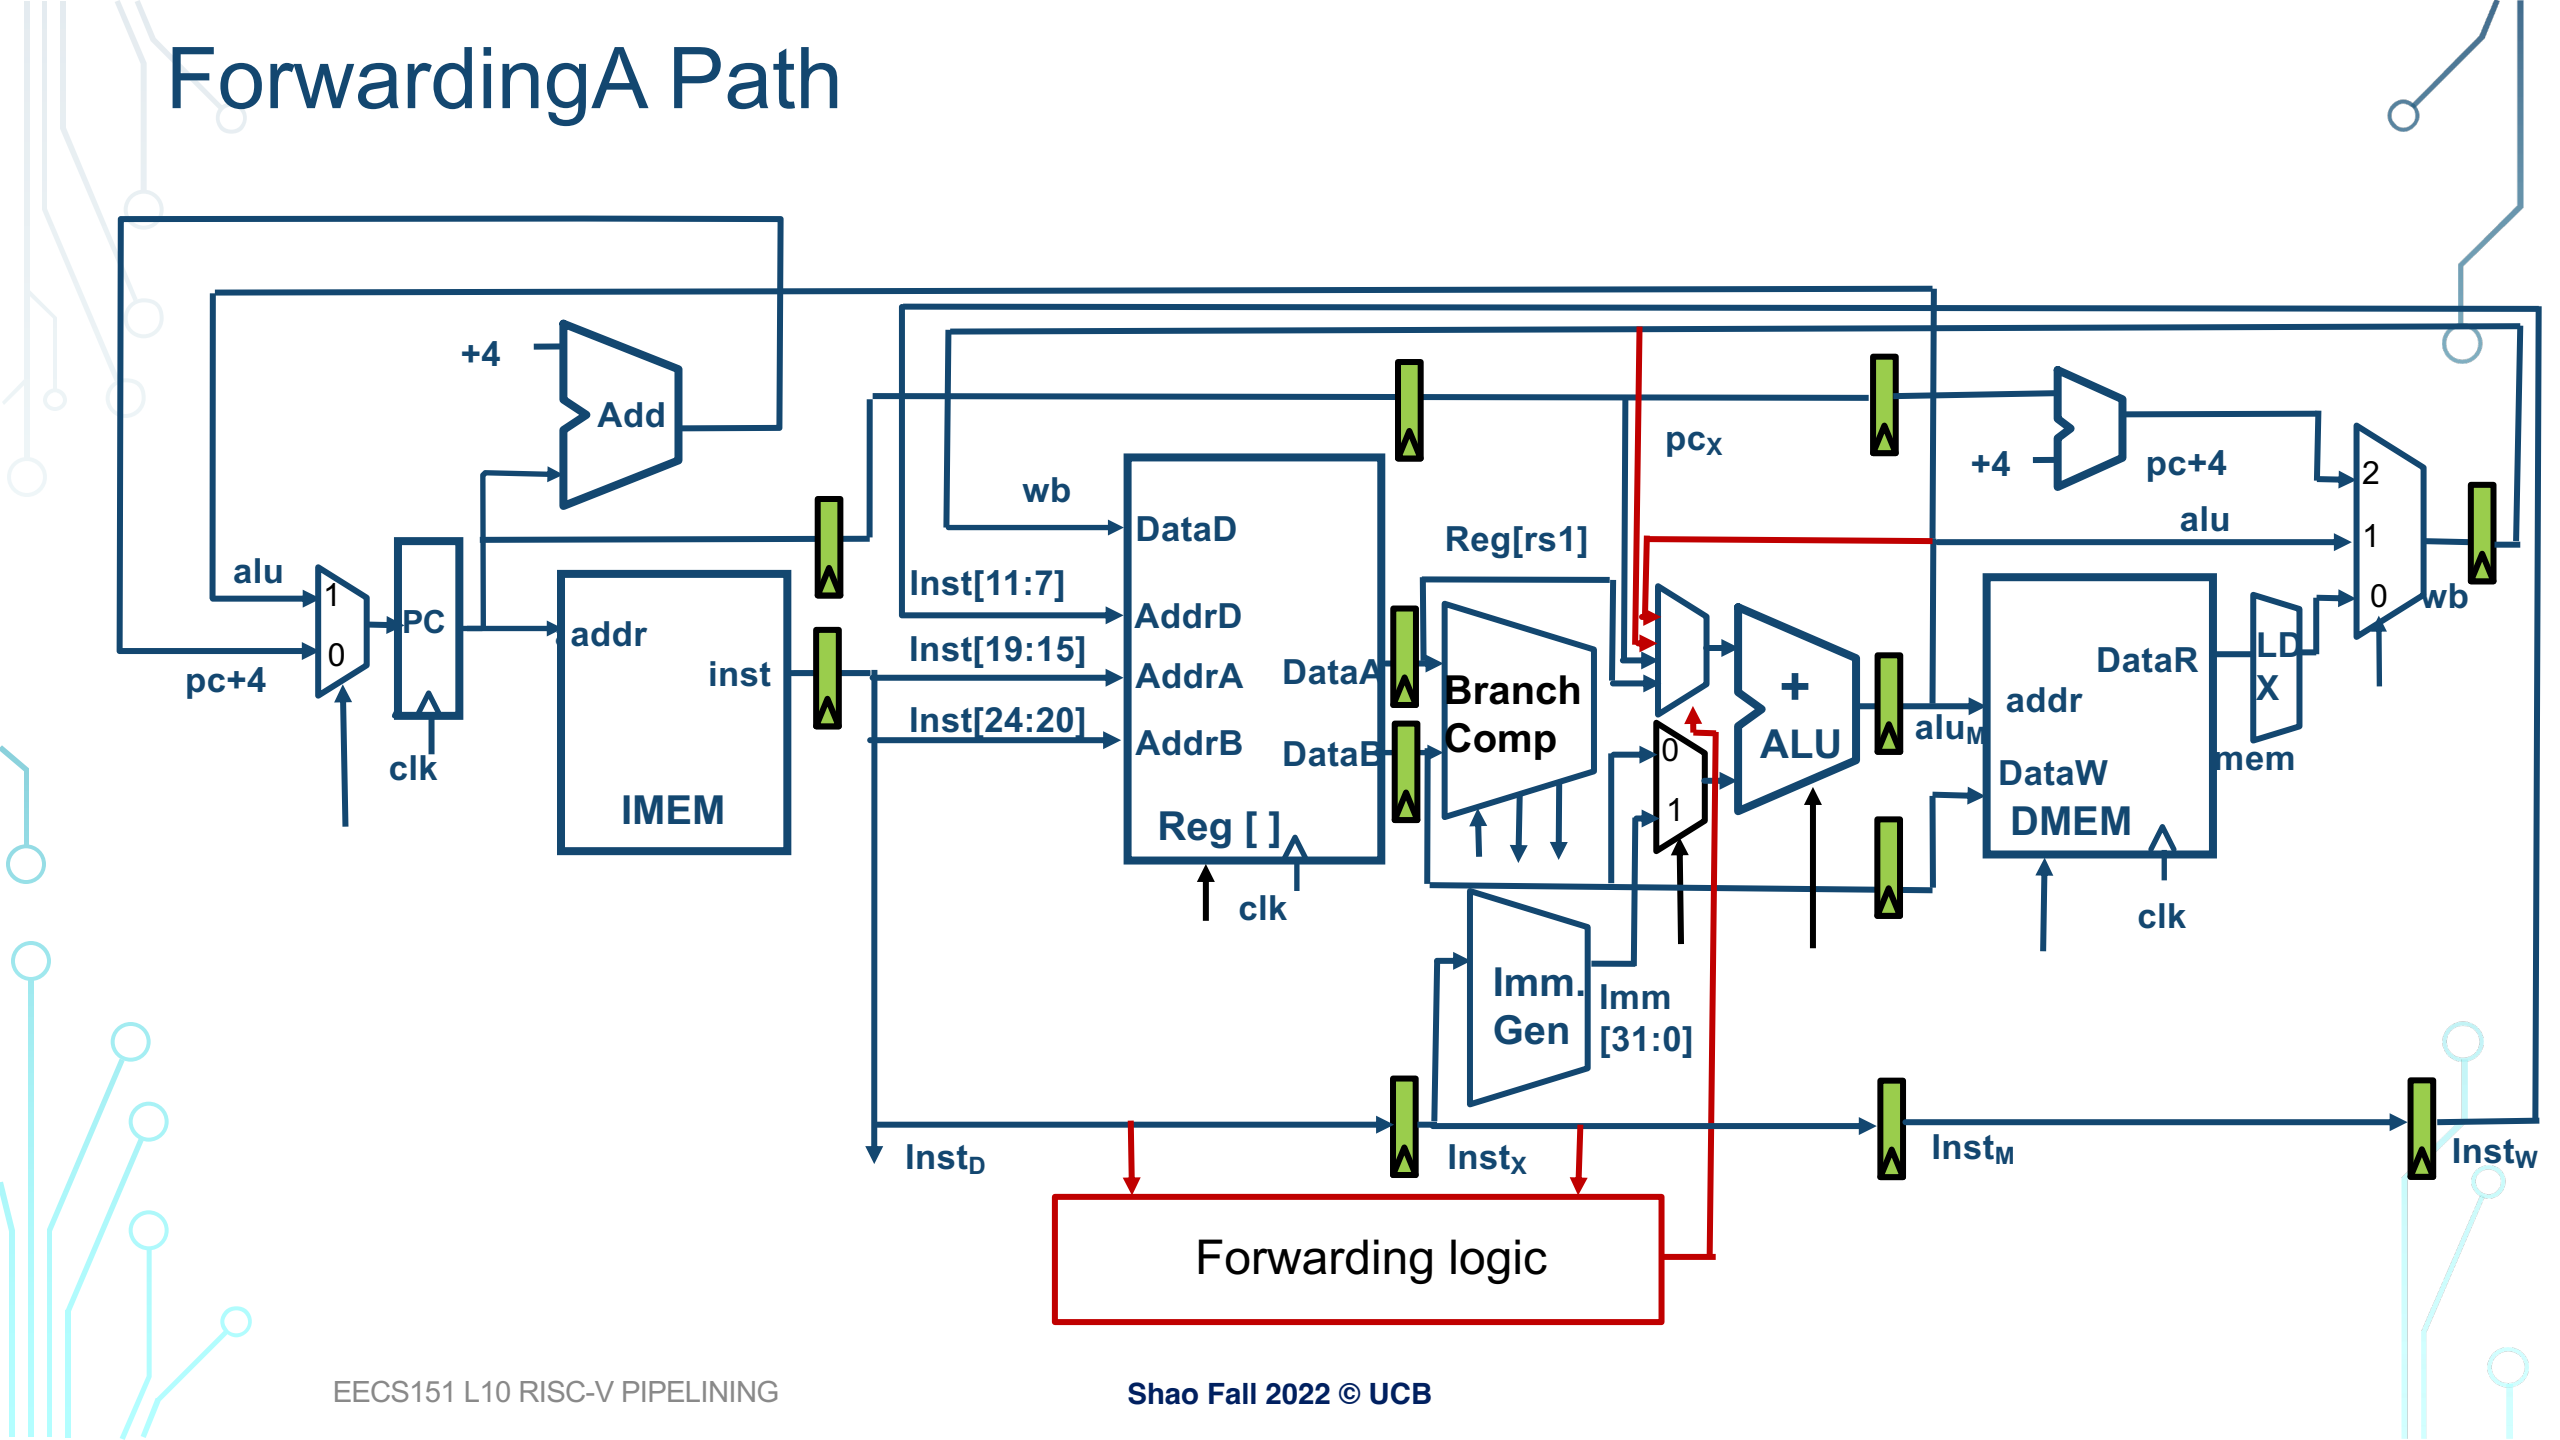
*Compare destination of older instructions in pipeline with sources of new instruction in decode stage.
Must ignore writes to x0!*

Forwarding Path

The diagram illustrates the forwarding path in a 5-stage RISC-V processor. The stages are Instruction (Inst), Decode (Dec), Execute (Ex), Memory (Mem), and Writeback (Wb). The Forwarding logic is shown as a red box at the bottom, which receives inputs from the Instruction, Decode, and Execute stages. It outputs a signal to the ALU in the Execute stage, which is used to forward the result of a previous instruction. The ALU also receives inputs from the Register File (Reg[rs1]) and the Immediate Generator (Imm. Gen). The ALU output is used to calculate the next PC value (pc+4) and to write back to the Register File. The Register File is updated with the result of the ALU operation. The Memory stage uses the ALU output to calculate the address for the Data Memory (DMEM). The Writeback stage uses the ALU output to write back to the Register File. The diagram shows the flow of data and control signals between these components, including the use of multiplexers and adders to calculate the next PC value and to forward the result of a previous instruction.

EECS151 L10 RISC-V PIPELINING

Shao Fall 2022 © UCB



Administrivia


- Midterm: 3/1, in class
 - Covers material up to this week.
 - One double-sided page of hand-written notes
 - RISC-V green card will be provided.
 - Review session next Tuesday 5-7pm.
 - Cory 540
- No new lab this week.
- New homework.



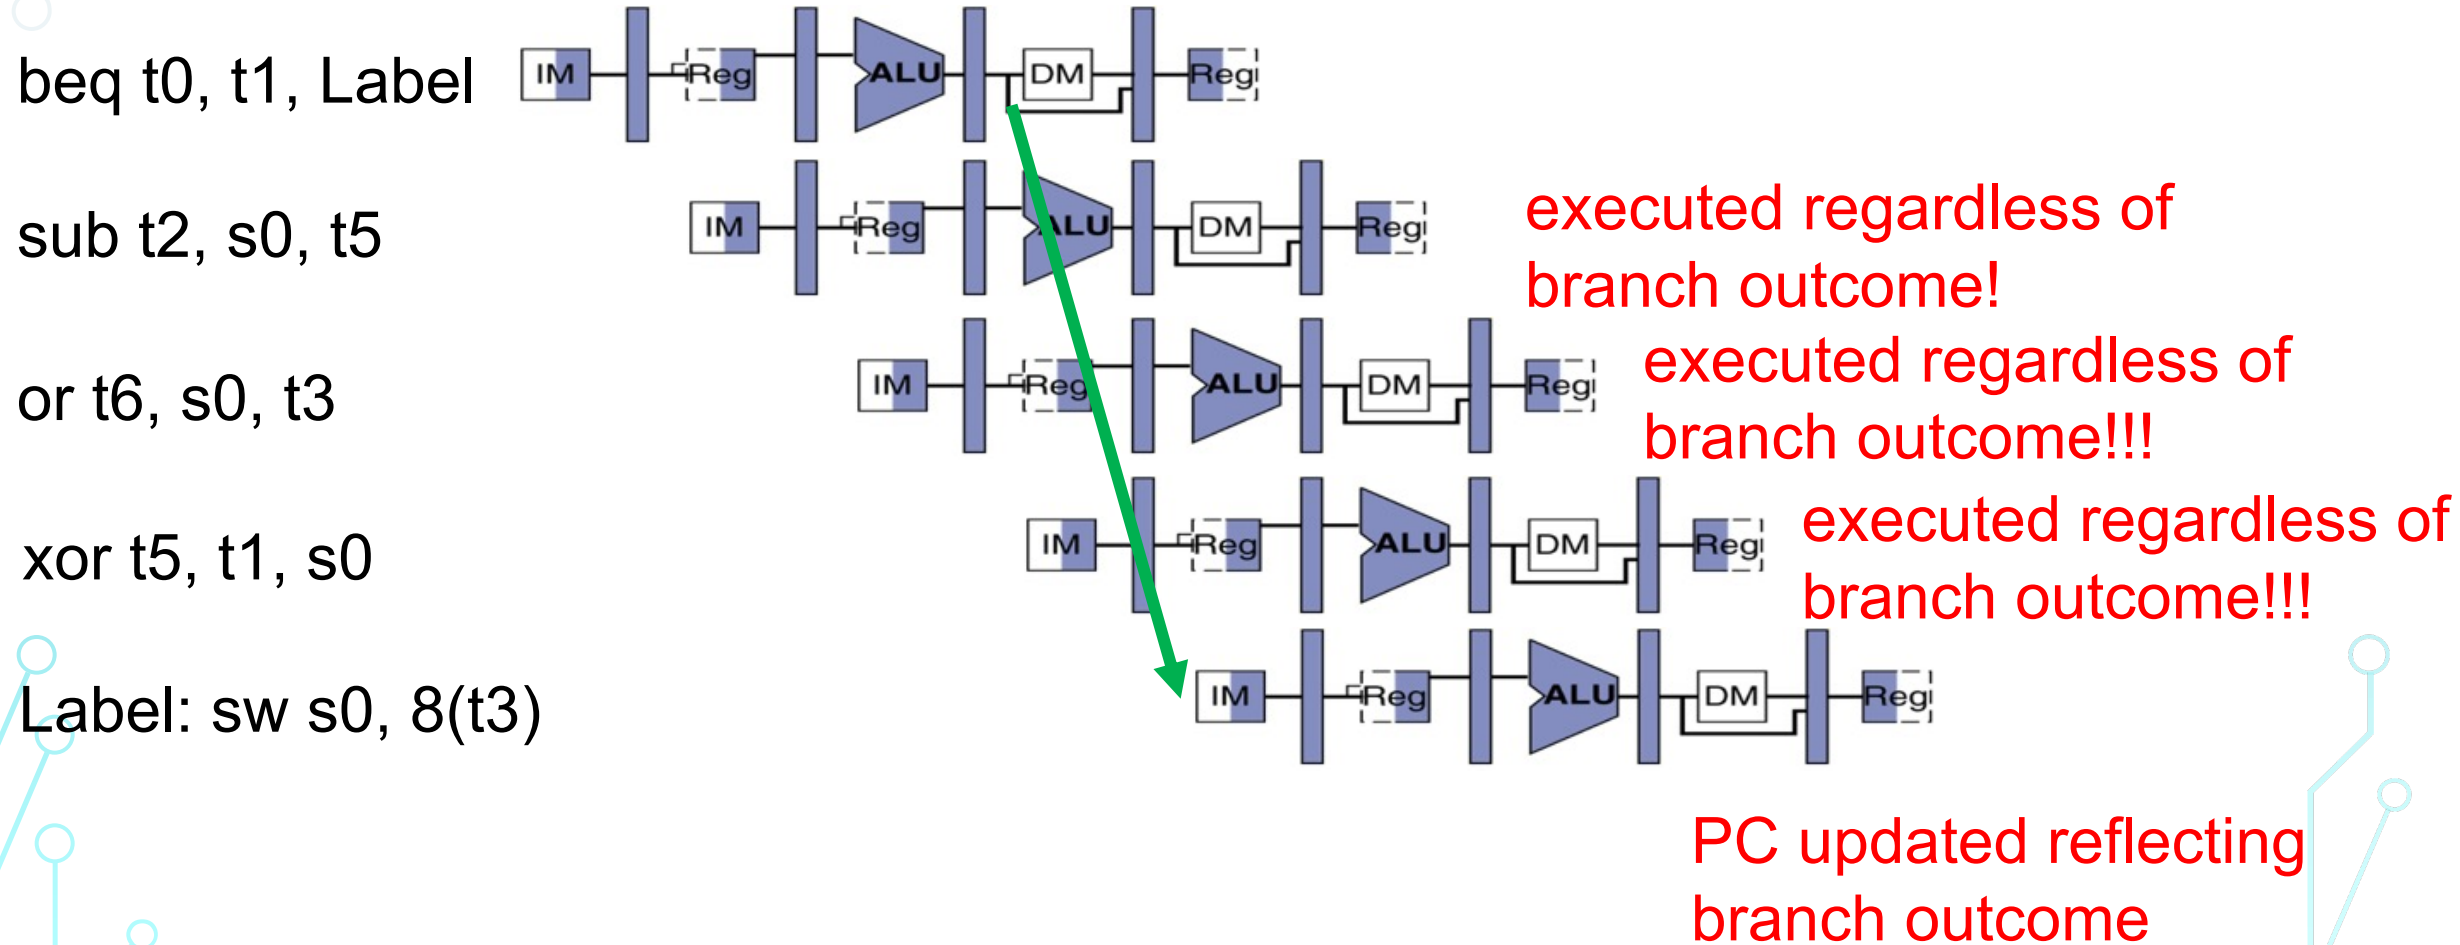
- **RISC-V Pipelining**

- **5-Stage Pipeline**
- **Pipeline Hazards**
 - **Structural**
 - **Data**
 - **Control**

Pipelined Datapath: When is the target PC calculated?



Control Hazards



Observation

- If branch not taken, then instructions fetched sequentially after branch are correct
- If branch or jump taken, then need to flush incorrect instructions from pipeline by converting to NOPs

Kill Instructions after Branch if Taken


beq t0, t1, label

sub t2, s0, t5

or t6, s0, t3

Add t2, t3, s0

label: xxxxxx



Taken branch

Convert to NOP

Convert to NOP

Convert to NOP

PC updated
reflecting branch
outcome

Reducing Branch Penalties

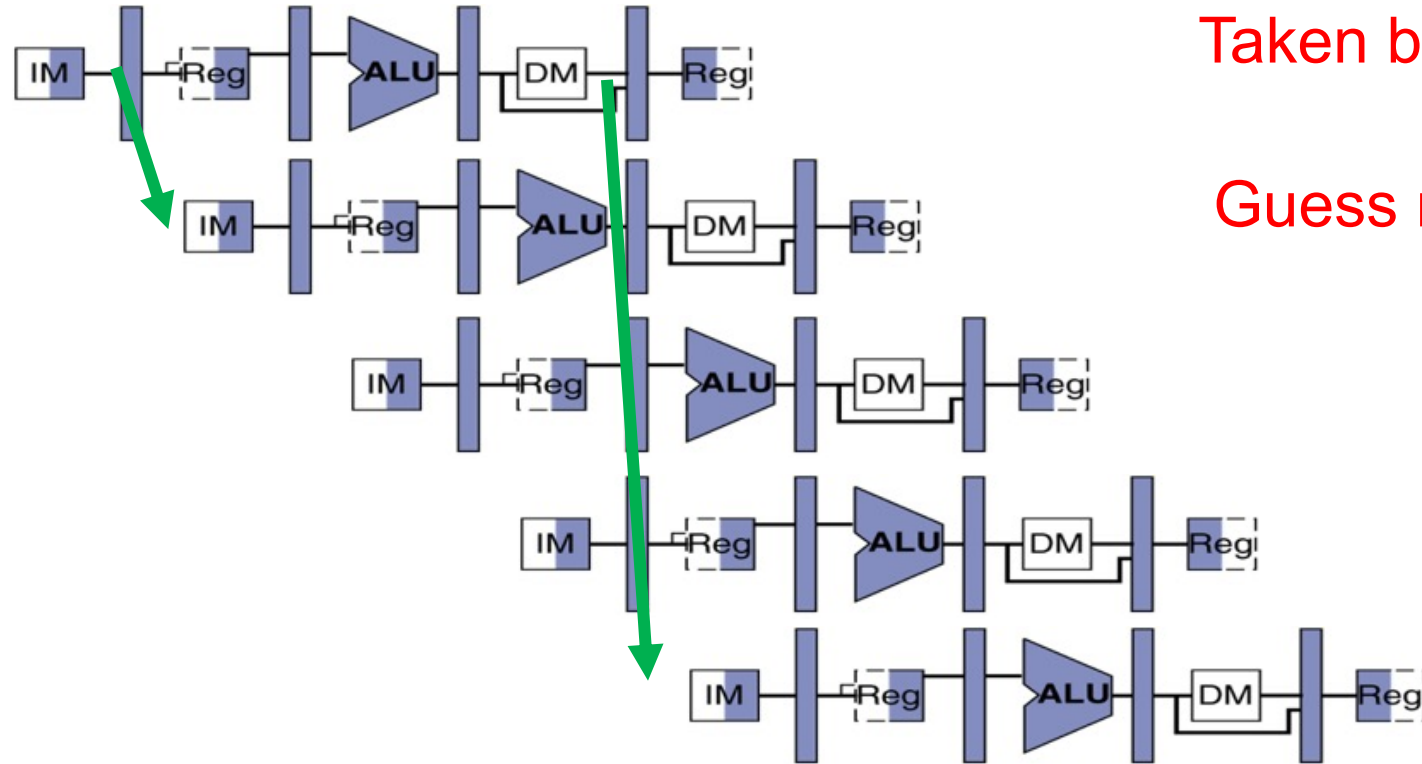
- In our datapath, every taken branch in simple pipeline costs 3 dead cycles
- To improve performance, use “branch prediction” to guess which way branch will go earlier in pipeline
- Only flush pipeline if branch prediction was incorrect

Branch Prediction

beq t0, t1, label

label:

.....



Taken branch

Guess next PC!

Check guess correct

Summary

- We have covered single-cycle, multi-cycle and pipelined datapaths
- Pipeline Hazards:
 - Structural
 - Data
 - Control hazards
- Need to be understood and handled appropriately
 - More resources
 - Forwarding
 - Stall